LLM ENGINEERING

Bonus chapter 2

Bonus: Long-Context Positional Extension

Self-paced reading

Interactive version: https://llm.strueber.dev/#/bonus-02-long-context

Bethgelab · University of Tübingen · Winter term 2026/27

Text and figures under CC BY 4.0, code under the MIT licence.

The problem

1

RoPE has no maximum position.

Why not serve an 8K-trained model at 32K?

Answer A1

Each pair turns at its own rate

(1) $$\omega_j = b^{-2j/d_r}, \qquad \text{angle at position } m = m\,\omega_j, \qquad \Delta = n - m$$

- Extension ratio $\rho = L_t / L_0$: target over trained length
- Fast pairs: short wavelength, many turns
- Slow pairs: long wavelength, few turns
- Every method below reshapes these bands

▼ **Pairs, frequencies and the extension ratio**

Formula (1) gives the rate at which each pair turns. The d_r rotated coordinates of a head form $d_r/2$ pairs, so d_r is even. The pairs are indexed $j = 0, \ldots, d_r/2 - 1$, and each rotates as a two-dimensional vector. At position m, pair j has the angle $m\omega_j$. The base b is not itself a frequency. It sets every pair's frequency through this formula. Code calls ω_j the inverse frequency because it is the reciprocal of b^{2j/d_r}.

Week 3 showed that the rotary part of the score between a query at position m and a key at position n depends on the two positions only through the signed offset $\Delta = n - m$. In causal attention a query sees only keys at or before its own position, so $n \le m$ and $\Delta \le 0$. A range of offsets is therefore written as a signed interval, such as $-(L_0 - 1) \le \Delta \le 0$ within training, or as a bound on $|\Delta|$.

The extension ratio $\rho = L_t / L_0 \ge 1$ divides the target length by the trained length. Serving an 8K-trained model at 32K means $\rho = 4$.

Write D_h for the head dimension. With full rotary, $d_r = D_h$ and every coordinate of the head rotates. With partial rotary, $D_h - d_r$ coordinates stay unrotated. An attention logit sums over the coordinates of the head, so with partial rotary it splits into a part from the rotated coordinates and a part from the unrotated ones.

The frequency falls as j grows. Low-index pairs rotate quickly: their wavelength is short, the same offset changes their angle a lot, and they complete many turns within the training length. High-index pairs rotate slowly: their wavelength is long, the same offset changes their angle little, and they complete few turns. Every method in this chapter changes how these bands behave when $L_t > L_0$.

PI and NTK

3

PI divides every angle by ρ

(2) $$m \mapsto \frac{m}{\rho}, \qquad \Delta \mapsto \frac{\Delta}{\rho}, \qquad \omega_j \mapsto \frac{\omega_j}{\rho}$$

- **Position Interpolation** (PI) (Chen et al., 2023)
- Adjacent angle: the rotation one extra token adds
- Position 32,767 maps to 8,191.75
- Fast pairs lose resolution too (Peng et al., 2024)

▼ **Rescaling positions rescales every frequency**

Position Interpolation (PI) makes a long input look like a short one (Chen et al., 2023). It divides every position by ρ, so position m becomes m/ρ and every offset Δ becomes Δ/ρ (Formula 2). A 32K input then spans the same range of positions that an 8K input spanned in training.

Claim. Dividing every position by ρ has the same effect as dividing every frequency by ρ. Under PI the query is at m/ρ and the key at n/ρ, so pair j sees the signed offset $(n - m)/\rho = \Delta/\rho$, and its relative angle is no longer $\Delta\,\omega_j$. The relative angle and the angle at a single position regroup in the same way:

$$\begin{aligned} (\Delta/\rho)\,\omega_j &= \Delta\,(\omega_j/\rho) && \text{relative angle at offset } \Delta/\rho \\ (m/\rho)\,\omega_j &= m\,(\omega_j/\rho). && \text{angle at position } m/\rho \end{aligned}$$

Rescaling the positions is therefore the same as rescaling the frequencies, and PI can be written as $\omega_j \mapsto \omega_j/\rho$, the third map in Formula (2). We call ω_j/ρ the PI frequency of pair j.

The **adjacent angle** of pair j is the angle that one extra token adds to its rotation, which is ω_j itself. To compare a remapping with the original, divide the new adjacent angle ω'_j by the old one. This adjacent-angle ratio ω'_j/ω_j is $1/\rho$ for every pair under PI: each pair turns ρ times more slowly per token than it did in training.

Take the illustrative values $L_0 = 8{,}192$, $L_t = 32{,}768$ and $\rho = 4$. The last target index, 32,767, maps to $32{,}767/4 = 8{,}191.75$. That value lies inside the half-open range $[0, 8{,}192)$ that the trained length spans, but just above 8,191, the largest integer index training contained. Offsets behave the same way: the largest target offset magnitude, $|\Delta| = 32{,}767$, maps to 8,191.75, three quarters of a position beyond the largest trained magnitude, 8,191. We use the nominal ratio $\rho = L_t/L_0$, not the ratio $(L_t - 1)/(L_0 - 1)$ that would align the integer endpoints exactly.

Chen et al. derive a stability bound for interpolation. In their LLaMA setup, they fine-tune each remapped model briefly, then check loss, retrieval, long-document tasks and quality within the original window (Chen et al., 2023).

Dividing every adjacent angle by ρ has a cost. The fast pairs change their angle a lot from one token to the next, and the model uses them to tell nearby positions apart. PI shrinks their adjacent angle by the same factor as the slow pairs', so they lose resolution too. This diagnosis in terms of frequencies comes from the later YaRN analysis, not from PI's original theorem (Peng et al., 2024).

Static NTK-aware scaling raises the base

$$b' = b\,\rho^{d_r/(d_r-2)} \tag{3}$$

- Named after the neural tangent kernel (NTK)
- Static: one fixed ρ sets one new base (bloc97, 2023)
- Dynamic NTK: recompute the base as length grows
- Adjusted base frequency (ABF): train with a larger base

▼ Three ways to change the base

A second family of methods leaves the positions alone and changes the base. Static NTK-aware scaling, from 2023, takes its name from the neural tangent kernel (NTK), through an argument we do not derive. It picks one fixed target ratio ρ and replaces the base b by $b' = b\,\rho^{d_r/(d_r-2)}$ (Formula 3) (bloc97, 2023). The formula first appeared in a community post, and maintained libraries implement it.

Claim. Static NTK-aware scaling multiplies the frequency of pair j by $\rho^{-2j/(d_r-2)}$. To see this, put b' into Formula (1):

$$\begin{aligned}
\omega'_j &= (b')^{-2j/d_r} && \text{Formula (1) with base } b' \\
&= b^{-2j/d_r}\,\rho^{-2j/(d_r-2)} && \text{Formula (3), power of a product} \\
&= \omega_j\,\rho^{-2j/(d_r-2)}. && \text{Formula (1)}
\end{aligned}$$

The adjacent-angle ratio is therefore $\omega'_j/\omega_j = \rho^{-2j/(d_r-2)}$. Its exponent is 0 at the fastest pair and -1 at the slowest pair, $j = d_r/2-1 = (d_r-2)/2$:

$$\begin{aligned}
\omega'_0/\omega_0 &= \rho^{0} = 1 && \text{fastest pair, } j=0 \\
\omega'_{d_r/2-1}/\omega_{d_r/2-1} &= \rho^{-1} = 1/\rho. && \text{slowest pair, } j = d_r/2-1
\end{aligned}$$

So the fastest pair keeps its angle, and the slowest pair gets the ratio $1/\rho$, the same as under PI. Between the two ends, each pair's ratio is the previous pair's ratio multiplied by the same factor $\rho^{-2/(d_r-2)}$, which is below 1 when $\rho > 1$. For the illustrative case $\rho = 4$ and $d_r = 64$, the new base is $b' = b \cdot 4^{64/62}$, about 4.183 times the old one.

Static scaling fixes ρ and b' once, so it applies its full remapping at every length, short inputs included: a 1K prompt gets the same remapped frequencies as a 32K one. Dynamic NTK recomputes a base from the observed sequence length as that length grows, so it remaps only

as far as the current input needs. The recomputation differs between variants, and Hugging Face documents the variants side by side (Hugging Face, n.d.).

Recomputing the base during generation constrains the KV cache: the cached keys were rotated under an earlier base, so the implementation must recompute or re-rotate them, or generate with one fixed base.

A trained high base, called adjusted base frequency (ABF), changes the same knob differently: you pick a larger b and train, or continue training, with it, so the weights adapt to the new base. ABF shares the knob with static NTK-aware scaling and Dynamic NTK, but not their algorithm.

Because all three methods change the base, a report that says only "the base was raised" does not say which method was used. A reproduction record names the variant, the original context, the target ratio, d_r, the base update policy, the cache policy and the training history.

5

What ratios do PI and NTK give?

$$\rho = 4,\ d_r = 64$$
$$\text{fastest pair } j = 0,\ \text{slowest pair } j = 31$$

Answer A2

Compare PI and NTK pair by pair

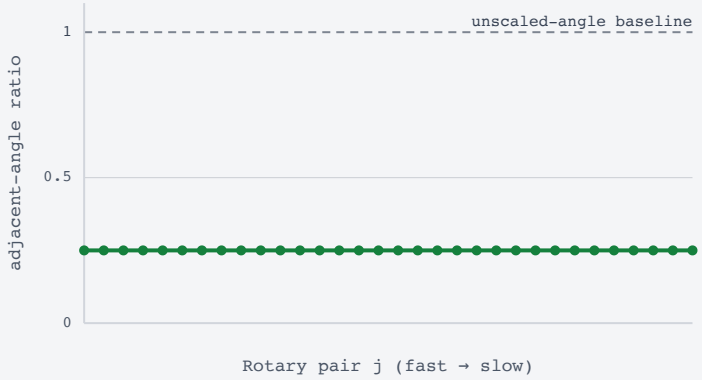
- Each dot: one rotary pair
- Height: new adjacent angle over the original

Adjacent-angle ratio

d_r: 32 64 ρ: 2x 4x 8x

method: No extension Position Interpolation NTK-aware

L_0=8K, L_t=32K, ρ=4: ratio by rotary pair



Select a rotary-pair dot to read its adjacent-angle ratio.

PI divides every adjacent angle by ρ. Static NTK-aware scaling leaves the fastest pair unchanged and moves the others step by step toward 1/ρ. Besides the No extension reference, the widget plots these two methods only. A ratio of 1 means a pair turns as in training; it says nothing about the length the model accepts or uses well. The curves are derived from the formulas; they are not measurements of model quality.

▼ Reading the adjacent-angle ratio

The widget plots the adjacent-angle ratio of every rotary pair. Each dot is one pair, from the fastest, $j = 0$, on the left to the slowest on the right. Its height is the pair's new adjacent angle divided by its original one, so a height of 1 means the pair turns as it did in training.

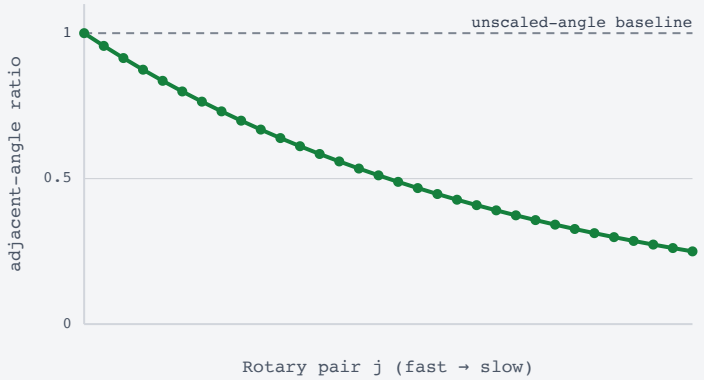
PI draws a flat line at $1/\rho$, and NTK-aware scaling falls from 1 to $1/\rho$. A larger ρ lowers PI's whole line and NTK-aware scaling's slow end, while NTK-aware scaling's fast end stays at 1. Changing d_r changes only the number of dots. The NTK-aware ratio $\rho^{-2j/(d_r-2)}$ equals $\rho^{-j/(d_r/2-1)}$. It depends only on how far along the axis the pair is, $j/(d_r/2-1)$, so the curve keeps its shape whatever d_r is.

Adjacent-angle ratio

d_r: 32 **64** ρ: 2x **4x** 8x

method: No extension | Position Interpolation | **NTK-aware**

L_0=8K, L_t=32K, ρ=4: ratio by rotary pair



Select a rotary-pair dot to read its adjacent-angle ratio.

PI divides every adjacent angle by ρ. Static NTK-aware scaling leaves the fastest pair unchanged and moves the others step by step toward 1/ρ. Besides the No extension reference, the widget plots these two methods only. A ratio of 1 means a pair turns as in training; it says nothing about the length the model accepts or uses well. The curves are derived from the formulas; they are not measurements of model quality.

NTK-aware scaling at ρ = 4: the fastest pair keeps a ratio of 1, and the ratio falls to 0.25 at the slowest pair.

YaRN and LongRoPE

7

YaRN blends by rotation count

$$\lambda_j = \frac{2\pi}{\omega_j} = 2\pi b^{2j/d_r}, \qquad r_j = \frac{L_0}{\lambda_j} \tag{4}$$

$$\omega'_j = (1-\gamma_j)\frac{\omega_j}{\rho} + \gamma_j\,\omega_j \tag{5}$$

- YaRN (Peng et al., 2024)
- Few rotations: interpolate like PI
- Many rotations: keep the original frequency
- In between: blend, called NTK-by-parts

▼ From rotation counts to the blend

YaRN treats the frequency bands selectively (Peng et al., 2024). Its rule needs two quantities for each pair. The first is the wavelength $\lambda_j = 2\pi/\omega_j$. Because ω_j is measured in radians per token and one full turn is 2π radians, λ_j is measured in tokens per cycle: the number of tokens pair j takes for one turn. With $\omega_j = b^{-2j/d_r}$ this gives $\lambda_j = 2\pi\, b^{2j/d_r}$ (Formula 4), so the wavelength grows with j.

The second quantity is the number of rotations the pair completes over the original context, $r_j = L_0/\lambda_j$. This is the paper's nominal count, and we keep it. Counting only the span between the first and the last position would give $(L_0 - 1)/\lambda_j$.

A ramp turns the rotation count into a blending weight:

$$\gamma_j = \mathrm{clip}\left(\frac{r_j - \alpha}{\beta - \alpha},\, 0,\, 1\right)$$

The weight is 0 for a pair with fewer than α rotations, such as a slow pair with its long wavelength. It is 1 for a pair with more than β rotations, such as a fast pair, and it rises linearly between the two. For the Llama family the paper recommends $\alpha = 1$ and $\beta = 32$ (Peng et al., 2024). The authors' reference code calls the same two values `beta_slow` and `beta_fast` (Quesnelle et al., 2023). Both are tunable choices made for that family.

Formula (5) uses the weight to mix two frequencies. At $\gamma_j = 0$ it gives ω_j/ρ, the PI frequency of pair j. At $\gamma_j = 1$ it gives the original ω_j. For a weight between 0 and 1, the coefficients $1 - \gamma_j$ and γ_j are nonnegative and add up to 1, so ω'_j is a convex blend of the two frequencies and

lies between them. The slow pairs are therefore interpolated as under PI, the fast pairs keep their original frequency, and the band between them is blended. YaRN's paper calls this selective blend NTK-by-parts interpolation. It is the step between the earlier methods and YaRN, which adds an attention scaling to it.

8

Which pairs does YaRN keep?

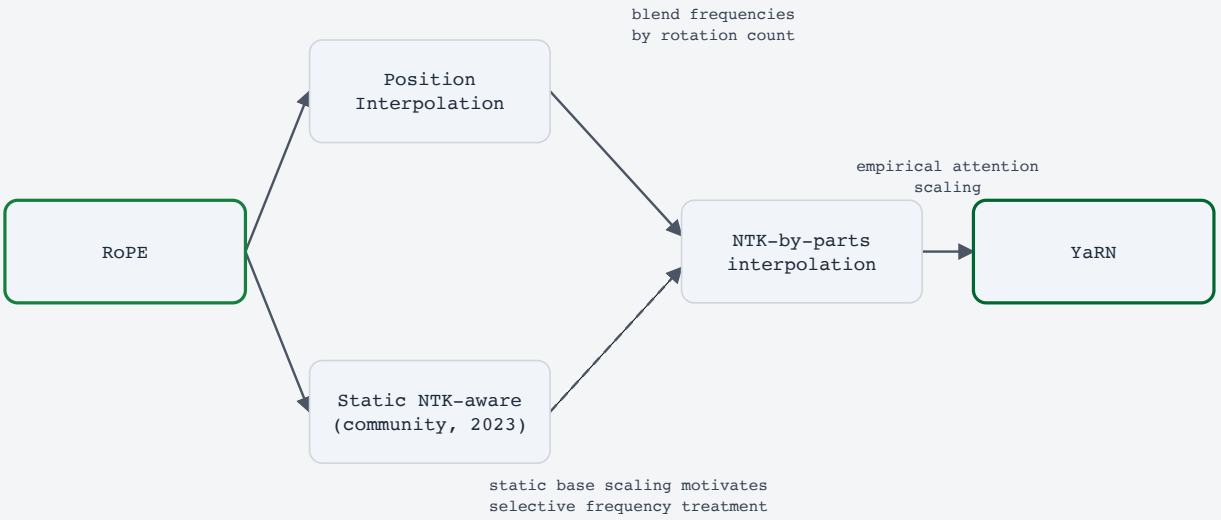
What is γ for pair j = 16?
L_0 = 8K, b = 10,000, d_r = 64
ramp α = 1, β = 32

Answer A3

YaRN adds attention scaling to the blend

$$a(\rho) = 0.1 \ln \rho + 1 \tag{6}$$

- Factor a fitted per ratio on Llama models (Peng et al., 2024)
- YaRN multiplies queries and keys by a
- Keeps attention entropy from growing with more keys



Schematic after Peng et al. (ICLR 2024), Figure 1.

A solid edge carries a mechanism into the next method; the dashed edge carries a motivating idea only. The edges show which idea each method takes from another, not that the methods are equivalent or that one inherits from another.

▼ Attention scaling and the method family

Remapping the frequencies is one half of YaRN. The other half is an attention scaling. Peng et al. fitted it on full-rotary Llama models, picking for each ratio the factor with the lowest perplexity (Peng et al., 2024). For the extension ratio ρ, the factor is $a(\rho) = 0.1 \ln \rho + 1$ (Formula 6), and YaRN multiplies both the query and the key by it. A logit is the dot product of a query and a key, so with full rotary every logit is multiplied by a^2. For $\rho > 1$ the factor exceeds 1, and it grows slowly with ρ.

Peng et al. link the factor to the increased entropy of attention at long inputs. As more keys enter each softmax, the attention weight spreads over more of them, and the entropy of the attention distribution grows. Multiplying every logit by a factor above 1 sharpens the softmax and lowers that entropy again. Check the formula and its constants on your own model before you reuse them. The authors' reference implementation shows how the ramp and the scaling are applied in code (Quesnelle et al., 2023).

The lineage figure, redrawn after Figure 1 of Peng et al., shows how YaRN grew out of the earlier methods (Peng et al., 2024). RoPE branches into Position Interpolation and static NTK-aware scaling, and both lead into NTK-by-parts interpolation. PI brings the mechanism, the interpolated frequencies that NTK-by-parts blends. Static NTK-aware scaling brings only the idea of treating the frequencies selectively, since it spares the fastest pair. The last edge, to YaRN, adds the attention scaling.

10

What is the new logit?

$\rho = 4$ gives $a \approx 1.1386$

$\ell_r = 2$ rotated, $\ell_u = 3$ unrotated

scale only the rotated Q and K

Answer A4

LongRoPE searches the rescale factors

- **LongRoPE** (Ding et al., 2024)
- Searched: one rescale factor per pair
- Searched: where rescaling starts
- Fine-tune to 256K, then search to 2,048K
- Factors readjusted for short contexts
- Used by the 128K Phi models (Abdin et al., 2024) (Microsoft et al., 2025)

▼ **Searching instead of a closed form**

PI, NTK-aware scaling and YaRN each compute the new frequencies from a closed-form rule. LongRoPE searches for them instead, guided by perplexity (Ding et al., 2024). It searches two non-uniformities. The first is one rescale factor for each rotary pair, so every pair can be slowed by its own amount. The second is a starting-position threshold: positions before it keep their original angles, and later positions get the searched rescaling. We do not derive the search.

The paper extends in stages. For LLaMA2 it fine-tunes first with factors searched for 128K and then with factors for 256K. A second search on the fine-tuned 256K model reaches 2,048K without further fine-tuning. Finally, it searches once more to readjust the factors for short contexts, such as 4K and 8K.

The schedule separates the length used to adapt the weights, at most 256K, from the much larger length reached by remapping, 2,048K. A low perplexity in the search or a nominal maximum length does not show that the model uses its context well. The evaluation matrix on the Evaluation tab decides that.

Among Microsoft's Phi models, only the long-context (128K) ones use LongRoPE. The Phi-3 report extends phi-3-mini to 128K with it, and the Phi-3.5 models use it to go from 4K to 128K (Abdin et al., 2024). Phi-4-mini reaches 128K with LongRoPE too (Microsoft et al., 2025).

Training

12

Three interventions, three costs

Intervention	Question it answers	Main cost
Inference-only remapping	Serve longer without training?	None up front; must pass the evaluation matrix
Continued long-context training	Can you afford to adapt the weights?	Training compute and long data
Target-length training	Must weights see the full length?	Attention, memory, batch and data

▼ Three interventions answer three questions

The remappings of the previous tabs change the angles. Reaching a longer context can also involve the weights, and three interventions answer three different questions.

Inference-only remapping asks whether you can serve longer inputs without any training. It changes the positional coordinates with one of the remappings above, such as PI, NTK-aware scaling or YaRN, and leaves the weights as they are. It costs nothing up front, so it is the cheapest intervention to try. It is acceptable only if the model passes the complete evaluation matrix of the Evaluation tab, against a threshold declared before the run.

Continued long-context training asks whether you can afford to adapt the weights. It trains the checkpoint further on longer sequences. This costs training compute and needs long training data, and in return the model gets a chance to learn to use distant evidence.

Target-length training asks whether the weights must see the full target length during training. In the chosen training stage it exposes them to sequences of the full target length, and it pays what such sequences cost: attention work, memory, a smaller batch and enough long data.

The three names describe what you do to the model. They do not tell you the compute, memory, data or quality cost that a given run ends up with. Compare those outcomes with the same evaluation matrix.

Real recipes mix training and remapping

Model	Training for length	Remapping
Llama 3 (Grattafiori et al., 2024)	Base 500,000 from the start; 405B: six stages, 8K to 128K	None in the report; Llama 3.1 code rescales by parts (Meta, 2025)
Qwen3 (Yang et al., 2025)	ABF to 1,000,000; a 32K stage	YaRN and Dual Chunk Attention (DCA) at inference
OLMo 3 (Ettinger et al., 2025)	8K to 65K on 50B to 100B tokens	YaRN on full-attention layers only
DeepSeek-V3 (DeepSeek-AI et al., 2024)	Two 1,000-step phases, 4K to 32K to 128K	YaRN on the RoPE part only
Long-context Phi (Abdin et al., 2024)	No length stage reported	LongRoPE, 4K to 128K

▼ What each recipe teaches

The table lists five recipes from model reports. Each row pairs what the reports say about training for length with the remapping the model uses, if any.

Llama 3 sets the RoPE base to 500,000 from the start of pre-training (Grattafiori et al., 2024). The report writes the base as θ, the course's b. For the 405B model the report then increases the context from the original 8K to 128K in six stages, so the weights see the full target length. The report mentions no YaRN. The report's models were released as Llama 3.1, and Meta's reference code for those weights rescales the frequencies by parts. The code uses the rotation count $r_j = L_0/\lambda_j$ of slide 7, over the original 8,192 tokens. Pairs with fewer than one rotation have their frequency divided by 8, pairs with more than four keep theirs, and the pairs between are blended (Meta, 2025). This is the mechanism of NTK-by-parts, with its own ramp boundaries, a factor of 8 instead of the extension ratio $L_t/L_0 = 16$, and no attention scaling.

Qwen3 raises its base from 10,000 to 1,000,000 with adjusted base frequency (ABF), the trained high base of slide 4, for a long-context pre-training stage at 32,768 tokens (Yang et al., 2025). At inference Qwen3 adds YaRN and Dual Chunk Attention (DCA), a method that splits a long input into chunks for attention. Together they let Qwen3 serve inputs four times longer than its 32K training length (Yang et al., 2025).

OLMo 3 extends its context from 8,192 to 65,536 tokens in a training stage of 50B tokens for the 7B model and 100B tokens for the 32B model (Ettinger et al., 2025). During that stage it applies YaRN only to its full-attention layers. Three of every four layers use sliding-window

attention, in which each token attends only to a fixed window of recent tokens. Those layers never see an offset beyond their window, however long the input, and OLMo 3 leaves their RoPE unchanged (Ettinger et al., 2025). Week 4 covers sliding windows.

DeepSeek-V3 extends from 4K to 32K and then to 128K in two training phases of 1,000 steps each, with YaRN applied in both (DeepSeek-AI et al., 2024). YaRN remaps only the RoPE part of its multi-head latent attention, the 64 rotated coordinates of the split described on slide 10.

The long-context Phi models reach 128K with LongRoPE, the searched rescaling of slide 11. Their reports describe no separate training stage for length. Phi-3 and Phi-3.5 credit the extension from 4K to 128K to LongRoPE (Abdin et al., 2024), and Phi-4-mini supports 128K based on it (Microsoft et al., 2025). LongRoPE as a method still fine-tunes at the lengths its search targets, as slide 11 describes.

The three YaRN rows differ in where and when they remap: in which layers, in which part of the attention and at which stage. Only Qwen3 applies YaRN at inference alone, while OLMo 3 and DeepSeek-V3 apply it during their extension training. No row is a default for every model.

Remapping pulls angles back; training improves use

- Remapping pulls angles back toward the trained range
- Remapping alone teaches nothing about distant evidence
- More extension tokens improved OLMo 3's long-context scores (Ettinger et al., 2025)
- PI's authors fine-tuned briefly after remapping (Chen et al., 2023)
- Longer sequences: more attention work and memory

▼ **Remapping and adaptation do different jobs**

A remapping changes only the position indices, or equivalently the frequencies, so that the angles of a long input move back toward the range training produced. How far each pair comes back depends on the method and the ratio. The model's internal states at the new lengths can still differ from anything it saw in training, and the remapping teaches it nothing about using distant tokens.

Adapting the weights on long data gives the model practice with distant tokens, and in the reports it improved long-context scores. OLMo 3 compared several token budgets for its extension stage. More tokens raised its scores on the long-context benchmark RULER, especially at longer sequence lengths (Ettinger et al., 2025). Position Interpolation, too, comes with adaptation: Chen et al. fine-tune briefly after interpolating (Chen et al., 2023).

Training on longer sequences also costs more. Longer sequences raise the attention work and the memory pressure of each step. The utilization a run actually reaches depends on the kernels, the parallelism, how sequences are packed and the hardware.

Evaluation

15

Context length is not context use

Check	Concrete 8K to 32K design
Language modeling	Negative log-likelihood (NLL) at 8K, 16K, 24K, 32K
Retrieval	Evidence at early, middle and late positions
Distant-evidence use	Question answering (QA), aggregation or summaries needing that evidence
Original window and serving	8K loss and tasks, short prompts, latency, memory
Robustness	Several prompts and documents; seeds or intervals
Reproducibility	Tokenizer, effective target, prompt format, decoding
Decision	Task floor and allowed regression, set in advance

▼ **Seven checks before calling an extension done**

Slide 12 asked every intervention to pass one evaluation matrix. The table makes its seven rows concrete for an extension from $L_0 = 8\text{K}$ to $L_t = 32\text{K}$, and a claim that the extended model works at 32K is credible only when it covers all seven.

Language modeling is measured as the negative log-likelihood (NLL) of held-out text: the average of $-\log p$ over the probabilities the model gives each actual next token. Report it at 8K, 16K, 24K and 32K tokens, so you can see whether prediction gets worse as the input grows. Its exponential, the perplexity, carries the same information and is optional.

Retrieval measures accuracy with the evidence an answer needs placed early, in the middle or late in the input. Slide 1 noted that failure depends on where the evidence appears, so each position is reported on its own.

Finding a string is one skill, and using what you found is another. The distant-evidence row therefore asks for at least one task whose answer requires the distant evidence: question answering (QA), aggregation or summarization over the long input.

The original-window and serving row rechecks the extended model where it already worked. Measure the loss and the tasks at 8K, the quality on short prompts, and the latency and memory of serving. A remapping can change the model's behavior on short inputs too: static NTK-aware scaling, for one, applies its full remapping at every length.

Robustness means several prompts and several documents, plus repeated seeds or confidence intervals, so that one prompt or one document does not decide the result. Reproducibility means recording the tokenizer, the effective target length the inputs reach in tokens, the prompt format and the decoding details.

The decision row states two numbers before the comparison runs: the floor the target task must reach, and the largest short-context or serving regression you will accept. A threshold chosen after the results are known can be fitted to them.

A nominal maximum context or a flat loss curve is not evidence of useful context on its own.

The ratios of PI, NTK-aware scaling and YaRN are algebra, and LongRoPE's factors come from a search. The attention-scaling constants, how well the model adapts on long data and its final quality can only be measured, and the matrix is where you measure them.

16

What is still missing?

8K to 32K: flat perplexity at 32K
retrieval passes early, middle and late

Answer A5

17

So why not serve an 8K-trained model at 32K?

Answer A6

Six results on long context

- Past L_0, slow pairs meet unseen angles
- PI scales all pairs, while NTK spares the fastest
- YaRN: blend by rotation count, scale attention
- LongRoPE: searched per-pair factors
- Recipes mix training and remapping
- Test context use, not context length

▼ Six results

Past the trained length L_0, RoPE still gives every position an angle, but the slow pairs reach angles that training never produced.

Position Interpolation divides the adjacent angle of every pair by ρ. Static NTK-aware scaling raises the base, keeps the fastest pair's angle and brings the slowest pair to $1/\rho$ of its angle. Dynamic NTK recomputes the base as the input grows, and ABF trains with a larger base.

YaRN keeps the frequencies of pairs with many rotations, interpolates pairs with few, blends the band between them and scales attention by a fitted factor.

Guided by perplexity, LongRoPE searches a rescale factor for each pair and the position where rescaling starts, and it extends in stages.

Inference-only remapping, continued long-context training and target-length training answer different questions at different costs. Most real recipes combine training on long data with a remapping or a high base: remapping pulls the angles back toward the trained range, and training improves the use of distant tokens.

Only an evaluation of context use against a threshold set in advance shows whether an extension works.

Reflection questions

No submission. Bring short answers to the Week 4 tutorial.

1. Take $b = 10{,}000$, $d_r = 64$ and an 8K-trained model served at 32K, so $\rho = 4$. Under static NTK-aware scaling, which of the seven pairs that never complete a turn within 8K still go beyond their largest trained angle at 32K, and by what factor does each exceed it? What do PI and YaRN do to the same seven pairs?
2. Your course model uses full-head RoPE with $b = 10{,}000$ and $D_h = 128$, trained on 1,024-token windows. It must now answer questions about 4,096-token documents, and you can afford one short continued-training run. Choose how to change the rotary frequencies and with which settings, say what the run trains on, and state the evaluation and the acceptance threshold you would fix before the run.
3. An 8K-trained model with full rotary, $b = 10{,}000$ and $d_r = 64$ must serve inputs of up to 32K tokens without any training, but most of its prompts are under 2K tokens. Compare PI and YaRN, both at $\rho = 4$: which pairs and which logits of a 1K prompt does each change, and by how much? Which would you deploy, and what must you measure before you decide?

References

20

RoPE and extension methods

- bloc97 (2023). NTK-Aware Scaled RoPE allows LLaMA models to have extended (8k+) context size without any fine-tuning and minimal perplexity degradation. Community post.
- Chen, S., Wong, S., Chen, L., & Tian, Y. (2023). Extending Context Window of Large Language Models via Positional Interpolation. arXiv preprint arXiv:2306.15595.
- Ding, Y., Zhang, L. L., Zhang, C., Xu, Y., Shang, N., Xu, J., Yang, F., & Yang, M. (2024). LongRoPE: Extending LLM Context Window Beyond 2 Million Tokens. *ICML*, PMLR 235, 11091-11104.
- Gelberg, Y., Eguchi, K., Akiba, T., & Cetin, E. (2026). Extending the Context of Pretrained LLMs by Dropping Their Positional Embedding. *ICLR*.
- Peng, B., Quesnelle, J., Fan, H., & Shippole, E. (2024). YaRN: Efficient Context Window Extension of Large Language Models. *ICLR*.
- Su, J., Ahmed, M., Lu, Y., Pan, S., Bo, W., & Liu, Y. (2024). RoFormer: Enhanced Transformer with Rotary Position Embedding. *Neurocomputing*, 568, 127063.

21

Model reports

- Abdin, M., Aneja, J., Awadalla, H., et al. (2024). Phi-3 Technical Report: A Highly Capable Language Model Locally on Your Phone. arXiv preprint arXiv:2404.14219.
- DeepSeek-AI, Liu, A., Feng, B., et al. (2024). DeepSeek-V3 Technical Report. arXiv preprint arXiv:2412.19437.
- Ettinger, A., Bertsch, A., Kuehl, B., et al. (2025). Olmo 3. arXiv preprint arXiv:2512.13961.
- Grattafiori, A., Dubey, A., Jauhri, A., et al. (2024). The Llama 3 Herd of Models. arXiv preprint arXiv:2407.21783.
- Merrill, W., Li, Y., Romero, T., et al. (2026). Olmo Hybrid: From Theory to Practice and Back. arXiv preprint arXiv:2604.03444.
- Microsoft, Abouelenin, A., Ashfaq, A., et al. (2025). Phi-4-Mini Technical Report: Compact yet Powerful Multimodal Language Models via Mixture-of-LoRAs. arXiv preprint arXiv:2503.01743.
- The Microsoft AI Team (2026). MAI-Thinking-1: Building a Hill-Climbing Machine. Technical report, Microsoft AI.
- Yang, A., Li, A., Yang, B., et al. (2025). Qwen3 Technical Report. arXiv preprint arXiv:2505.09388.

22

Implementation sources

- DeepSeek-AI. (2025). DeepSeek-V3 inference code (inference/model.py).
- Hugging Face. (n.d.). Utilities for Rotary Embedding.
- Meta. (2025). Llama 3 reference model code (models/llama3/model.py).
- Quesnelle, J., Peng, B., Fan, H., & Shippole, E. (2023). YaRN reference implementation.

23

Image credits

- Dürer, A. (1525). Tall graduated scale with a fan of lines and a small dial, a construction for proportional measure. *Underweysung der Messung*. SLUB Dresden. Public domain.

Answers

A1 · Slow pairs leave the trained range

- Illustrative: $b = 10{,}000$ and $d_r = 64$
- Every position still gets an angle
- Slow pairs reach angles never seen in training

	Within 8K	At 32K
Slowest pair's largest angle	1.1 rad	4.4 rad
Pairs below one full turn	7 of 32	2 of 32

▼ Why serving past the training length is a distribution shift

Suppose a RoPE-based checkpoint was trained with contexts up to $L_0 = 8\text{K}$ tokens and must now serve inputs of $L_t = 32\text{K}$ tokens, without retraining from scratch. RoPE has no maximum position, so why not serve the checkpoint at 32K as it is?

Here L_0 is the largest context represented in the relevant training stage. It can be smaller than the maximum position the configuration allows, so read it from the training report or records.

Recall from Week 3 that RoPE splits each head's d_r rotated coordinates into $d_r/2$ pairs. Pair j turns by $\omega_j = b^{-2j/d_r}$ radians per token, where b is RoPE's base. The first pairs turn fast and the last ones slowly. A pair's angle grows linearly with position and has no upper limit (Su et al., 2024), so the model accepts a 32K input. What changes is what it sees. Past L_0, the input holds absolute indices that the training stage never contained, and relative offsets of larger magnitude than any it contained.

Take the illustrative values $b = 10{,}000$ and $d_r = 64$, so each head has 32 rotary pairs. The fastest pair turns by 1 rad per token and completes about 1,300 full turns within 8K tokens, so training has shown it every angle many times. The slowest pair turns by only about $1.33 \cdot 10^{-4}$ rad per token. Its wavelength, the number of tokens one full turn takes, is about 47,100 tokens. Within 8K its angle reaches only about 1.1 rad. At 32K it reaches about 4.4 rad. Seven of the 32 pairs never complete one turn within 8K tokens. At 32K, 2 of the 32 still stay below one turn.

At 32K the slowest pairs take angles that training never produced. This is the argument YaRN makes for treating the slow pairs differently from the fast ones (Peng et al., 2024). Real models differ: a higher base, such as the 500,000 of Llama 3, leaves more pairs below one turn within 8K.

How badly a model fails on such inputs depends on the model, the task, the length and where the evidence appears in the input. So measure how quality changes with input length in your own setting instead of assuming a cliff at a fixed length.

This chapter is required reading. It has no submission, and later lectures do not depend on it. It ends with three reflection questions, which the Week 4 tutorial discusses, so read it before that tutorial and prepare a short answer to each. The chapter is about extending a RoPE-based checkpoint. NoPE and

hybrid positional designs, which Week 3 covers, are the alternative path. Some recent models avoid the problem by leaving RoPE out of their long-range layers. The global attention layers of MAI-Thinking-1 use no positional encoding (Microsoft AI Team, 2026). OLMo Hybrid removes RoPE for its long-context stage (Merrill et al., 2026) with DroPE, a method that drops the positional embeddings of a pretrained model (Gelberg et al., 2026).

Back to the question

A2 · PI shrinks every pair; NTK spares the fastest

Adjacent-angle ratio	PI	Static NTK-aware
Fastest pair, $j = 0$	0.25	1
Slowest pair, $j = 31$	0.25	0.25

- PI: one factor for every pair
- NTK: the fastest pair keeps its angle

▼ Two ratios at the two ends

With $\rho = 4$ and $d_r = 64$, PI divides every adjacent angle by 4. Its ratio ω'_j/ω_j is therefore 0.25 at every pair, from the fastest, $j = 0$, to the slowest, $j = 31$.

For static NTK-aware scaling, substitute into the ratio $\rho^{-2j/(d_r-2)}$, whose exponent here is $-2j/62$. At $j = 0$ the ratio is $4^0 = 1$. At $j = 31$ it is $4^{-62/62} = 0.25$. A pair in the middle shows the shape between the ends: at $j = 16$ the ratio is $4^{-32/62} \approx 0.489$, almost twice PI's 0.25.

Both methods bring the slowest pair to $1/\rho$. They differ in what they do to the fast pairs: PI slows them by the full factor, while NTK-aware scaling leaves the fastest pair unchanged and slows its neighbours only a little.

Back to the question

A3 · At 8K, YaRN keeps 13 pairs and interpolates 7

- Pairs 0–12: over 32 turns, kept
- Pairs 25–31: under one turn, interpolated
- Pairs 13–24: blended
- Pair $j = 16$: $r = 13.04$ and $\gamma = 0.388$

▼ Walking the ramp across 32 pairs

Take YaRN's ramp for the Llama family, $\alpha = 1$ and $\beta = 32$, with the illustrative values $L_0 = 8\text{K}$ (8,192 tokens), $b = 10{,}000$ and $d_r = 64$.

Claim. With these values, pair j completes $r_j = 8{,}192 \cdot 10{,}000^{-j/32}/(2\pi)$ rotations over the original context. Combine Formula (4) with $\omega_j = b^{-2j/d_r}$:

$$\begin{aligned} r_j &= \frac{L_0}{\lambda_j} && \text{Formula (4)} \\ &= \frac{L_0\,\omega_j}{2\pi} && \text{Formula (4)}, \lambda_j = 2\pi/\omega_j \\ &= \frac{8{,}192 \cdot 10{,}000^{-j/32}}{2\pi}. && \text{Formula (1), values above} \end{aligned}$$

Going from one pair to the next divides the count by $10{,}000^{1/32} \approx 1.33$. The fastest pair turns about 1,304 times. Pair 12 still turns about 41 times and pair 13 about 31 times, so pairs 0 to 12 have $r_j > 32$. Their weight is $\gamma_j = 1$, and they keep their original frequency. At the other end, pair 24 turns about 1.3 times and pair 25 about 0.98 times, so pairs 25 to 31 have $r_j < 1$. Their weight is $\gamma_j = 0$, and they are interpolated as under PI. These are the 7 pairs that slide 1 found never complete one turn within 8K tokens: with $\alpha = 1$, the ramp's lower boundary is one full turn. Pairs 13 to 24 fall between the two boundaries and are blended.

Claim. Pair $j = 16$ is blended with the weight $\gamma_{16} \approx 0.388$. Its frequency in radians per token, its wavelength in tokens per cycle and its rotation count follow from Formulas (1) and (4):

$$\begin{aligned} \omega_{16} &= 10{,}000^{-1/2} = 0.01 && \text{Formula (1)} \\ \lambda_{16} &= 2\pi/0.01 \approx 628.3 && \text{Formula (4)} \\ r_{16} &= 8{,}192/\lambda_{16} \approx 13.04 && \text{Formula (4)} \\ \gamma_{16} &\approx (13.04 - 1)/(32 - 1) \approx 0.388. && \text{ramp, since } 1 < r_{16} < 32 \end{aligned}$$

With $\rho = 4$, PI's frequency for this pair is $0.01/4 = 0.0025$, and Formula (5) gives

$$\begin{aligned}\omega'_{16} &\approx (1 - 0.388) \cdot 0.0025 + 0.388 \cdot 0.01 \quad \text{Formula (5)} \\ &\approx 0.00541.\end{aligned}$$

The blended frequency lies between PI's 0.0025 and the original 0.01.

Back to the question

A4 · Scaling only rotated coordinates gives 5.59

$$a^2\ell_r + \ell_u \approx 1.2965 \cdot 2 + 3 \approx 5.593$$

- Scaling the whole logit instead gives 6.482
- DeepSeek-V3's code scales the whole logit (DeepSeek-AI, 2025)

▼ **Partial rotary changes what scaling touches**

Write ℓ for a head's scaled logit, the query-key dot product divided by $\sqrt{D_h}$, and split it into a rotated and an unrotated part, $\ell = \ell_r + \ell_u$. The dot product of a query and a key is a sum over the head's D_h coordinates. The part ℓ_r collects the terms of the d_r rotated coordinates, ℓ_u collects those of the $D_h - d_r$ unrotated ones, and both carry the factor $1/\sqrt{D_h}$. Scaling one coordinate of the query and the same coordinate of the key by a multiplies their product by a^2.

Claim. Scaling all of Q and K by a multiplies the logit by a^2, and scaling only the rotated coordinates multiplies only ℓ_r by a^2:

$$\ell' = \begin{cases} a^2\ell & \text{scale every coordinate,} \\ a^2\ell_r + \ell_u & \text{scale only the rotated coordinates.} \end{cases}$$

With full rotary there are no unrotated coordinates, so $\ell_u = 0$ and the two lines agree: scaling Q and K by a gives $\ell' = a^2\ell$. With partial rotary they differ, and scaling only the rotated coordinates gives $\ell' = a^2\ell_r + \ell_u$.

Claim. For $\rho = 4$ and the illustrative parts $\ell_r = 2$ and $\ell_u = 3$, scaling only the rotated part gives about 5.593, and scaling the whole logit gives about 6.482:

$$\begin{aligned} a = 0.1 \ln 4 + 1 &\approx 1.1386 && \text{Formula (6), } \rho = 4 \\ a^2 &\approx 1.2965 && \text{square} \\ a^2\ell_r + \ell_u &\approx 1.2965 \cdot 2 + 3 \approx 5.593 && \text{only the rotated part} \\ a^2(\ell_r + \ell_u) &\approx 1.2965 \cdot (2 + 3) \approx 6.482. && \text{the whole logit} \end{aligned}$$

The two results come from two implementation choices. Before you reuse a YaRN implementation, check which part it scales.

DeepSeek-V3 makes the whole-logit choice. It uses multi-head latent attention (MLA), which Week 4 covers. In MLA, each head's query and key have a 128-coordinate part without RoPE and a 64-coordinate part that carries RoPE, and the key's RoPE part is shared by all heads. Here $d_r = 64$ of $D_h = 192$ coordinates rotate. YaRN's frequency remap acts only on the RoPE part (DeepSeek-AI et al., 2024). Yet when the maximum sequence length exceeds the original, the

reference code multiplies the softmax scale $1/\sqrt{D_h}$ by a^2, so it scales every logit (DeepSeek-AI, 2025). The part without RoPE is scaled too.

Back to the question

A5 · Flat perplexity and retrieval are not enough

- A task that needs distant evidence
- Quality at 8K, plus serving latency and memory
- Several prompts, documents and seeds
- Tokenizer and decoding details recorded
- A threshold set before the run

▼ What the two passing checks leave open

Take an extension from 8K to 32K whose perplexity stays flat up to 32K and whose retrieval passes with the evidence early, in the middle and late. These two checks are the first two rows of the matrix. Flat perplexity says that the model predicts the next token about as well at 32K as at shorter lengths. Passing retrieval says that it can find a string anywhere in the input. Neither shows that the model can use what it finds.

The other five rows of the evaluation matrix are still open: a task that needs the distant evidence, such as aggregating facts from both ends of a document; the quality at 8K and the latency and memory of serving at 32K; several prompts, documents and seeds; a record of the tokenizer, the prompt format and the decoding; and a threshold chosen before the run.

Back to the question

A6 · Remap the angles, train for use, measure both

- Every position still gets an angle
- Slow pairs leave the trained range
- Remapping pulls the angles back toward the trained range
- Training on long data improves use of distant tokens
- Only the evaluation matrix decides

▼ The answer in five steps

So why not serve an 8K-trained model at 32K? First, the model accepts the input, because RoPE gives every position an angle.

Second, the slow pairs then reach angles that the training stage never produced, and the model meets inputs unlike any it was trained on.

Third, a remapping changes the frequencies to pull the angles back toward the trained range. PI, static NTK-aware scaling, YaRN and LongRoPE each change the pairs differently.

Fourth, a remapping changes only the coordinates. In the reports, training on long data improved the model's use of distant tokens. Four of the five recipes on slide 13 train the weights for length. The exception is Phi, whose reports describe no separate training stage for length. Every row adds a remapping, a high base, or both.

Fifth, only the evaluation matrix shows whether an extension worked, measured on the extended model against a threshold set in advance.

Back to the question